Apprentice

Show it once. Delegate forever.

The agent that learns your issue triage by watching you work, proves it with backtests and regression tests, then earns autonomy on a trust ladder.

Same issue. Same four clicks. Again.

Bug reports keep arriving with no version and no steps. Triage is the same ritual each time: label it, ask for a repro, assign it, next. Agents could help, but today that means writing a prompt and then hoping it's right.

GITHUB.COM/BUILDERIO/AGENT-NATIVE #8301

Error screen: Something went wrong

version: none · steps: none

GITHUB.COM/BUILDERIO/AGENT-NATIVE #8141

Error screen: Something went wrong

version: none · steps: none

THE RITUAL

+ needs-info

+ bug

request repro

assign → triage

REAL OPEN ISSUES IMPORTED READ-ONLY FROM THE AGENT-NATIVE REPO

A click and an agent call are the same typed action.

In Agent-Native, the UI and the agent share one set of actions. So every triage click is already a structured demonstration: action name, typed arguments, the issue it touched. No screen scraping, no guessing.

- Maya clicks “+ needs-info”
- The agent calls label-issue
- Same action → recorded as a demonstration

```
// actions/label-issue.ts
export default defineAction({
  schema: z.object({
    issueId, add, remove,
    playbookContext,
  }),
  ...domainAction("label-issue", …),
});
```

Watch. Prove. Then act.

1 WATCH

Typed clicks become demonstrations

Every action Maya takes is recorded with its arguments and the issue's features.

2 NOTICE

Two similar ones become a pattern

A ghost chip predicts the next step. A pattern card proposes a playbook.

3 PROVE

Backtest, shadow, tests

Replay on history, review in shadow mode, and every correction becomes a regression test.

4 EARN

Autonomy on a trust ladder

Watch → Suggest → Draft → Act → Auto, gated by the framework's approval hook. Undo any run.

It learns from what she already does.

After two hand-triaged issues, Apprentice predicts the next step (Tab to accept) and proposes a playbook. It even notices teammate Bob does it differently and lets the team choose.

- Pattern found · seen 2×
- Your team handles this two ways

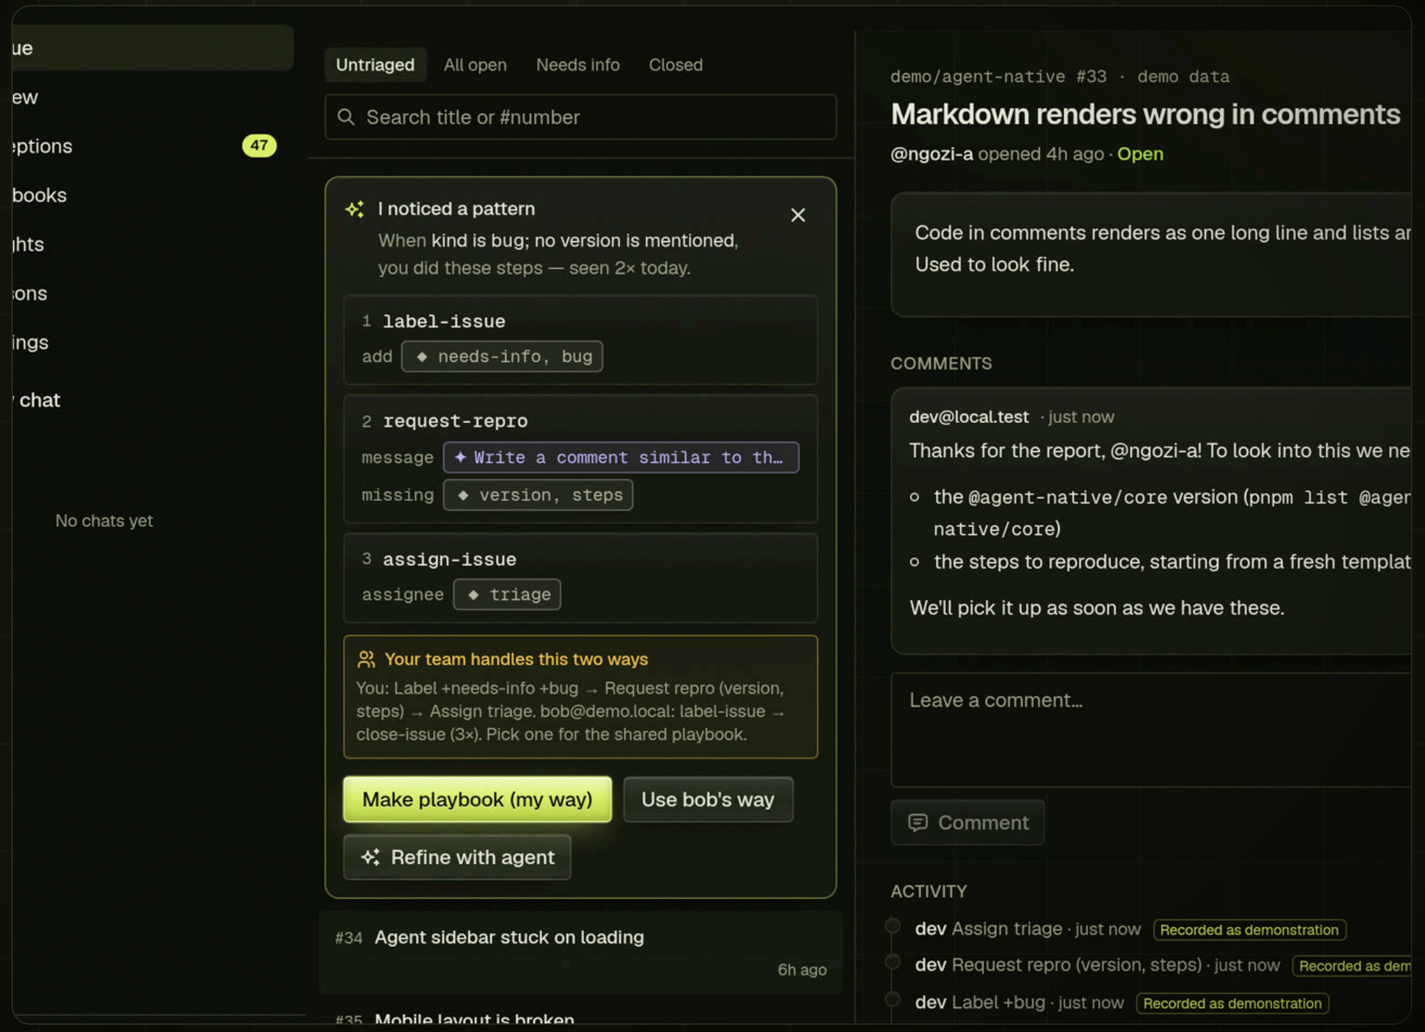

DEMO DATA · LOCAL APP · ISSUE #33

Numbers, not vibes.

CONSIDERED	MATCHED	AGREED	FALSE POSITIVES
29	17	14	3

Toggle a condition and the numbers update live against her own triage history. Loosening “has a version” raised matches to 21, but false positives jumped from 3 to 7. So she kept it tight.

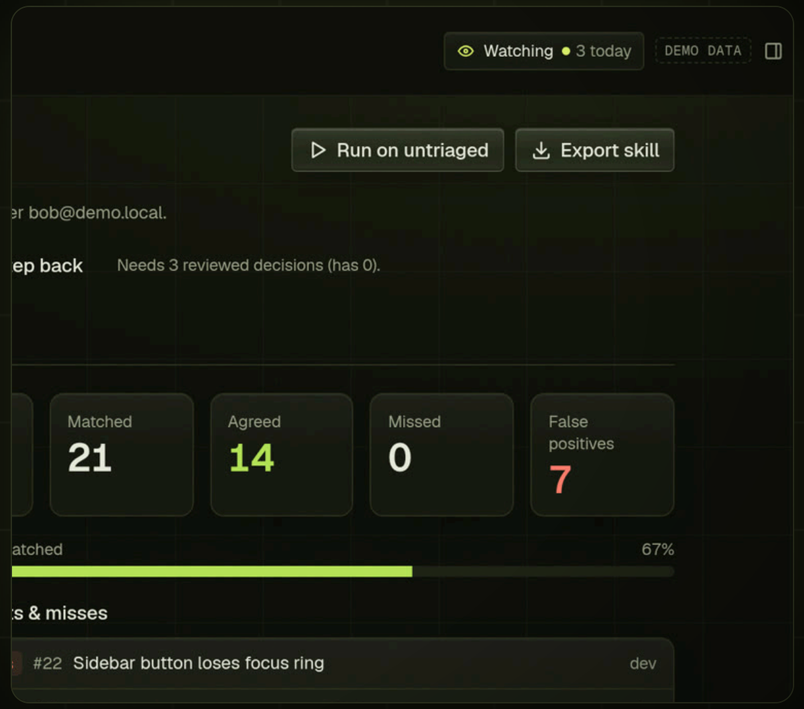

LIVE BACKTEST · AFTER LOOSENING

Every promotion needs evidence and a human click.

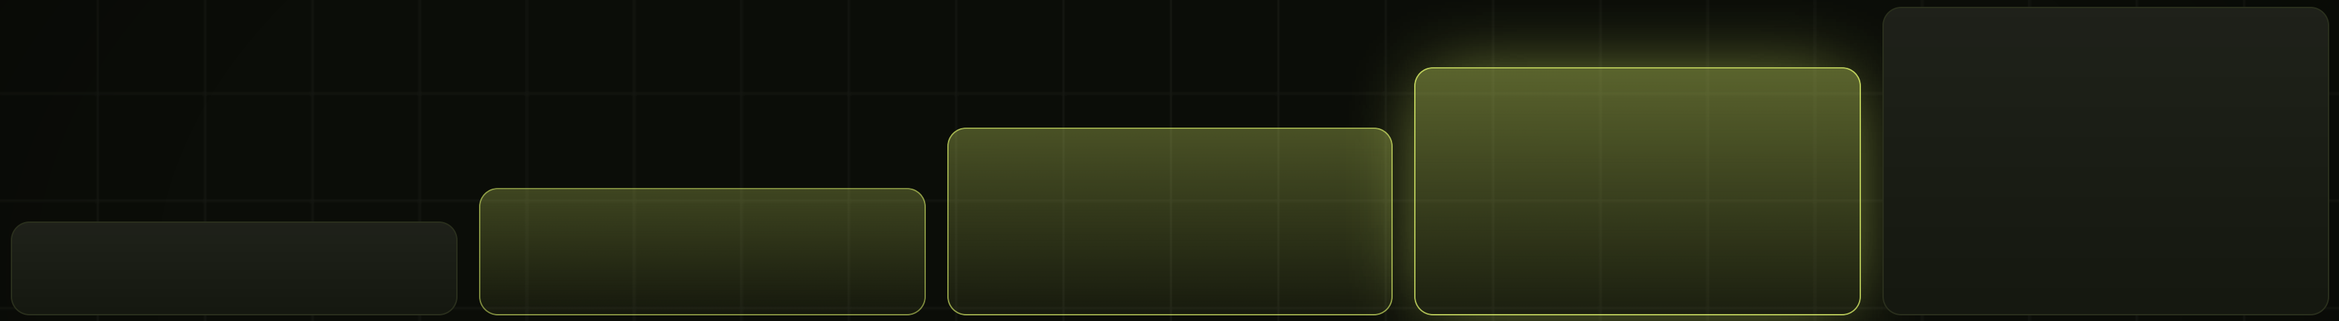

PROMOTION	REVIEWED	AGREEMENT	EXECUTED	ALWAYS
Suggest → Draft	≥ 3	≥ 80%	—	every regression test passing + a human click. An override at Act or Auto steps it back one rung.
Draft → Act	≥ 5	≥ 85%	—	
Act → Auto	≥ 8	≥ 90%	≥ 5	

One correction becomes a test.

Real issue #8918 is a feature request, not a bug, so the rule got it wrong. Maya overrides it. Apprentice saves the correction as regression test #3 and steps the playbook back from Act to Draft until it passes.

- Saved as test #3
- Stepped back to Draft: override at autonomy

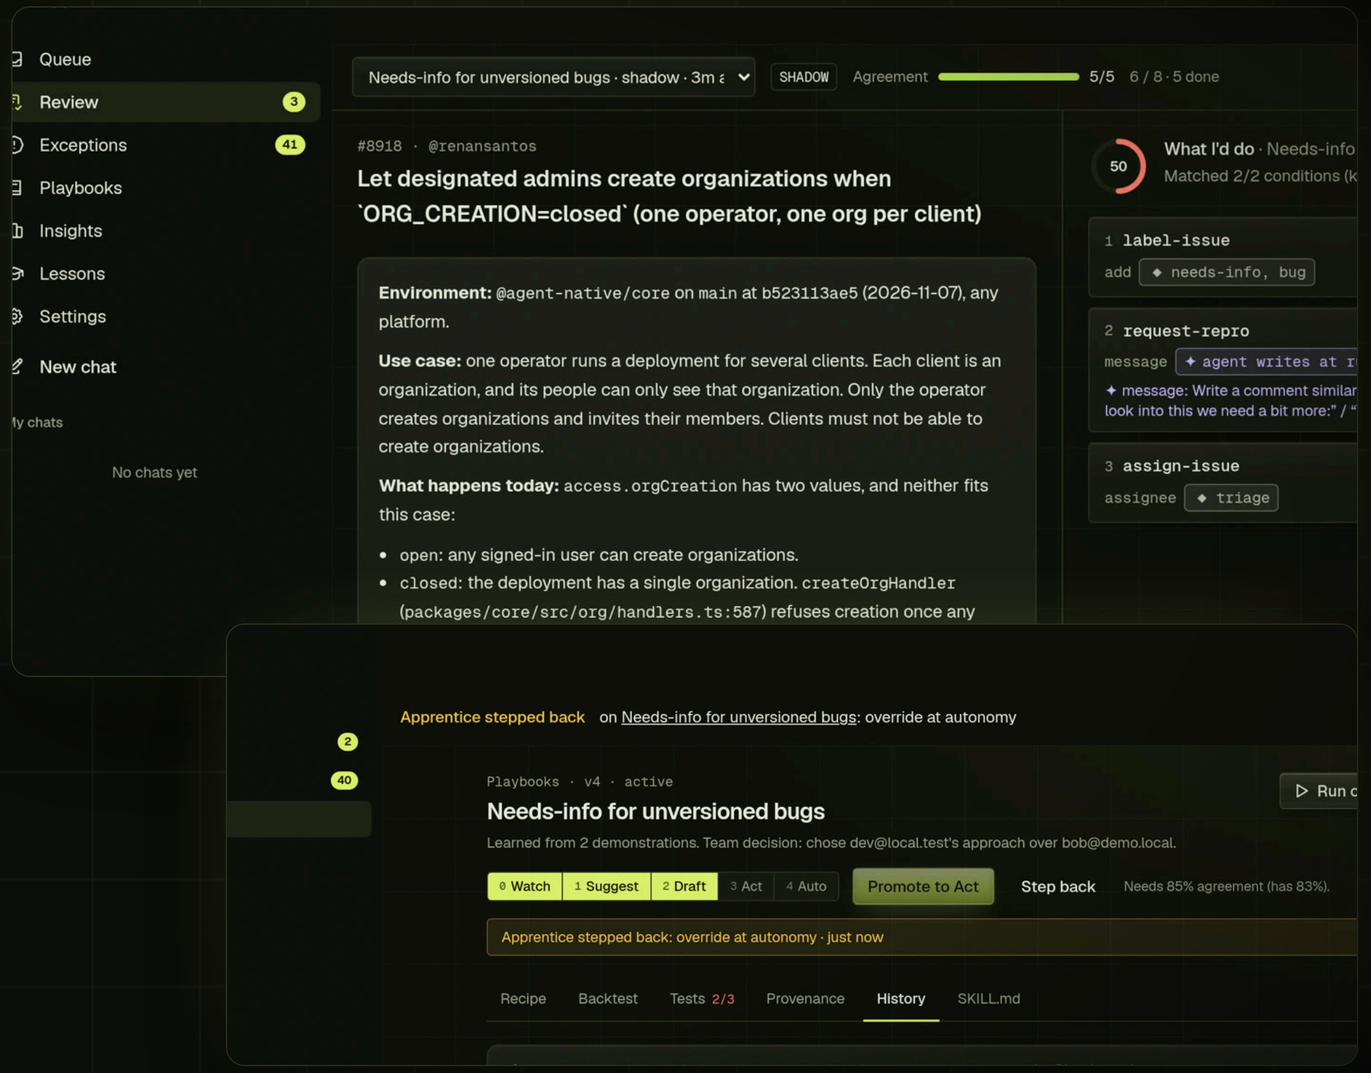

REAL ISSUE #8918 · OVERRIDE → REGRESSION TEST

“Who taught you this?” It shows you.

Every action Apprentice takes links back to the conditions that matched and the exact demonstrations it learned from. At Act, each run is handled on its own and comes with one-click undo.

- Taught by her: #31, #32
- 4 handled · each with undo

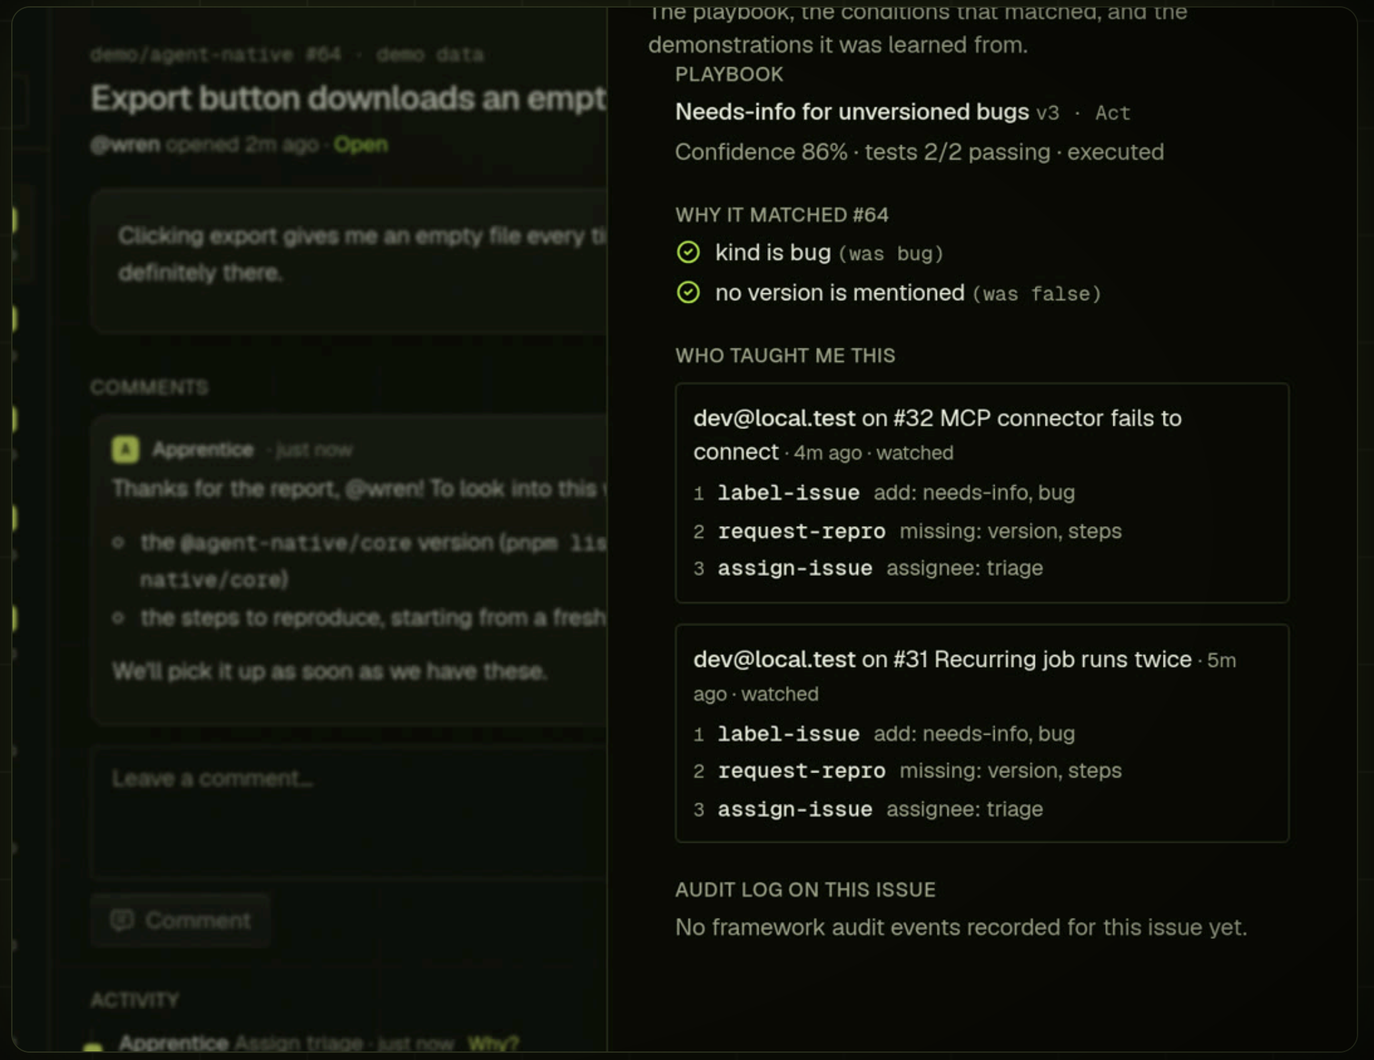

The framework is the product.

AGENT-NATIVE	WHAT APPRENTICE DOES WITH IT
`defineAction`	62 typed actions shared by the UI, the agent and MCP. Clicks are demonstrations.
`needsApproval`	The trust ladder: playbook steps ask before acting below Act.
`view-screen`	The agent sees the queue, issue, playbook or decision you're on.
Skills	triage-domain, apprentice-teacher, playbook-runner, shadow-reviewer.
MCP + automations	Review batches from Claude or ChatGPT; hourly autopilot + digest.
Export	Every playbook exports as a portable SKILL.md any agent can use.

```
// services/approval.ts
if (pb.trustLevel <= LEVELS.DRAFT)
  return true;
if (pb.trustLevel === LEVELS.ACT)
  return !!step?.checkpoint
    || destructive;
return false;
```

She only sees what needs her.

HANDLED BY APPRENTICE

4

EST. TIME SAVED

11m

AGREEMENT WITH YOU

83%

Autopilot runs the active playbooks on new issues and posts a digest. Exceptions are grouped by why they still need a human. Time saved is an estimate from her own median handling time.

Real, running, tested.

62

typed actions

55

tests passing, incl. the full learn → prove → act → undo loop

30

real GitHub issues imported in the demo

JUDGING CRITERION	WHERE TO LOOK
Usefulness & execution · 40%	Complete flow: teach → backtest → shadow → promote → act → explain → undo, on real issues
Creativity & design · 20%	Ghost-chip prediction, live backtest toggles, trust-ladder UI, keyboard-first triage
Innovation · 20%	Learning from typed demonstrations; corrections become regression tests
Use of Agent-Native · 20%	Shared actions, needsApproval as the trust ladder, view-screen, skills, MCP, automations

Apprentice

Show it once.
Delegate forever.

APPRENTICE-AGENT.ONRENDER.COM

APPRENTICE-AGENT.PAGES.DEV

GITHUB.COM/LSUDOKO/APPRENTICE

BUILT WITH **AGENT-NATIVE** · BUILDER.IO AGENT-NATIVE HACKATHON 2026